

Practical 1 Report

Task 1

Task 1.1

SELECT COUNT(*) AS "NUMBER OF PLAYERS" FROM ATHLETE2 WHERE CCODE = 'AUS';
<code>SELECT COUNT(*) AS "NUMBER OF PLAYERS" FROM ATHLETE2 WHERE CCODE = 'AUS';</code>
NUMBER OF PLAYERS ----- 243

Task 1.2

SELECT SPORTID AS "SportID", COUNT(*) AS "Count" FROM ATHLETE3 WHERE CCODE = 'RUS' GROUP BY SPORTID;	
<code>SELECT SPORTID AS "SportID", COUNT(*) AS "Count" FROM ATHLETE3 WHERE CCODE = 'RUS' GROUP BY SPORTID;</code>	
SportID	Count
-----	-----
30	5
51	1
44	2
29	6
47	4
53	11
46	23
52	17
50	39
45	4
49	25
11 rows selected	

Task 1.3

CREATE TABLE "ATHLETE_FULL" ("ATHLETEID" NUMBER, "FNAME" VARCHAR2(30 BYTE), "SNAME" VARCHAR2(30 BYTE), "BDATE" VARCHAR2(30 BYTE), "CCODE" VARCHAR2(30 BYTE), "SPORTID" NUMBER);
INSERT INTO "ATHLETE_FULL" SELECT * FROM ATHLETE1;
INSERT INTO "ATHLETE_FULL" SELECT * FROM ATHLETE2;
INSERT INTO "ATHLETE_FULL" SELECT * FROM ATHLETE3;
DESC "ATHLETE_FULL";

```
SELECT COUNT(*) FROM ATHLETE_FULL, COUNTRY WHERE ATHLETE_FULL.CCODE =
COUNTRY.CCODE AND (CONTINENT = 'EU');
```

```
DROP TABLE ATHLETE_FULL;
```

```
CREATE TABLE "ATHLETE_FULL"
("ATHLETEID" NUMBER,
"FNAME" VARCHAR2(30 BYTE),
"SNAME" VARCHAR2(30 BYTE),
"BDATE" VARCHAR2(30 BYTE),
"CCODE" VARCHAR2(30 BYTE),
"SPORTID" NUMBER);
```

```
INSERT INTO "ATHLETE_FULL" SELECT * FROM ATHLETE1;
INSERT INTO "ATHLETE_FULL" SELECT * FROM ATHLETE2;
INSERT INTO "ATHLETE_FULL" SELECT * FROM ATHLETE3;
```

```
DESC "ATHLETE_FULL";
```

```
SELECT COUNT(*) FROM ATHLETE_FULL, COUNTRY WHERE ATHLETE_FULL.CCODE = COUNTRY.CCODE AND (CONTINENT = 'EU');
```

```
DROP TABLE ATHLETE_FULL;
```

Query 1) Creating the ATHLETE_FULL table.
table "ATHLETE_FULL" created.

Query 2) Inserting the rows from ATHLETE1, ATHLETE2, ATHLETE3 respectively.

```
7,655 rows inserted.
9,662 rows inserted.
7,274 rows inserted.
```

Query 3) Displaying the attributes of ATHLETE_FULL table.

```
DESC "ATHLETE_FULL"
Name      Null Type
-----
ATHLETEID  NUMBER
FNAME      VARCHAR2(30)
SNAME      VARCHAR2(30)
BDATE      VARCHAR2(30)
CCODE      VARCHAR2(30)
SPORTID    NUMBER
```

Query 5) Dropping the ATHLETE_FULL table to avoid duplication in the count.

```
table ATHLETE_FULL dropped.
```

Query 4) Demonstrating the count of the athletes from Europe.

```
COUNT(*)
-----
16077
```

NOTE: The ATHLETE_FULL table is dropped at the end of the query due to the count being duplicated multiple times when the script is run. This is because every time the script is re-run the same tables (ATHLETE1, ATHLETE2, ATHLETE3) will be re-inserted into ATHLETE_FULL, thereby duplicating the count of athletes whose continent is Europe. It is important to drop ATHLETE_FULL so to ensure the accurate count.

Task 2

For all the following queries, BEGIN and END commands needed to be removed and only the nested query (Line 2-4 in all users below) in the user script called s4556162 was run to demonstrate the output screenshotted. Otherwise, the following appears and does not specify whether the ATHLETEID = 305 row, has been updated:

anonymous block completed

This is because the nested/executable part of the query is written within a code block. When the message “anonymous block completed” shows up, it means that there are no errors within the nested query (between the BEGIN and END commands) and the code block has successfully executed. To confirm that the row where ATHLETEID = 305 has actually been updated, we must remove the BEGIN and END commands (Line 1, 5, 6 in all users below). The BEGIN and END commands were left in the code and screenshots below because this is the standard SQL practice to update a database table (Oracle, 2022).

Job 1 – Full Replication	
USER1_HF_FULL_S4556162	<pre>BEGIN UPDATE USER1_HF_FULL_S4556162.ATHLETE1_REPLICA1 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>
	<pre>BEGIN UPDATE USER1_HF_FULL_S4556162.ATHLETE1_REPLICA1 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>
	1 rows updated.
USER2_HF_FULL_S4556162	<pre>BEGIN UPDATE USER2_HF_FULL_S4556162.ATHLETE1_REPLICA2 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>
	<pre>BEGIN UPDATE USER2_HF_FULL_S4556162.ATHLETE1_REPLICA2 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>
	1 rows updated.
USER3_HF_FULL_S4556162	<pre>BEGIN UPDATE USER3_HF_FULL_S4556162.ATHLETE1_REPLICA3 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>

	<pre>BEGIN UPDATE USER3_HF_FULL_S4556162.ATHLETE1_REPLICA3 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>
	1 rows updated.
Job 2 – Partial Replication	
USER1_HF_PA_S4556162	<pre>BEGIN UPDATE USER1_HF_PA_S4556162.ATHLETE1_REPLICA1 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>
	<pre>BEGIN UPDATE USER1_HF_PA_S4556162.ATHLETE1_REPLICA1 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>
	1 rows updated.
USER2_HF_PA_S4556162	<pre>BEGIN UPDATE USER2_HF_PA_S4556162.ATHLETE1_REPLICA2 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>
	<pre>BEGIN UPDATE USER2_HF_PA_S4556162.ATHLETE1_REPLICA2 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>
	1 rows updated.
USER3_HF_PA_S4556162	<p>Pass – does not need to be updated as this database user only contains the following tables:</p> <ul style="list-style-type: none"> • ATHLETE2_REPLICA2 • ATHLETE3_REPLICA3
Job 3 – No Replication	
USER1_HF_NO_S4556162	<pre>BEGIN UPDATE USER1_HF_NO_S4556162.ATHLETE1_REPLICA1 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>
	<pre>BEGIN UPDATE USER1_HF_NO_S4556162.ATHLETE1_REPLICA1 SET CCODE = 'AUS' WHERE ATHLETEID = 305; END; /</pre>

	1 rows updated.
USER2_HF_NO_S4556162	Pass – does not need to be updated as this database user only contains the following tables: <ul style="list-style-type: none"> ATHLETE2_REPLICA2
USER3_HF_NO_S4556162	Pass – does not need to be updated as this database user only contains the following tables: <ul style="list-style-type: none"> ATHLETE3_REPLICA3

Task 3

SELECT FNAME, SNAME, BDATE FROM USER1_VF_S4556162.ATHLETE_V1, USER2_VF_S4556162.ATHLETE_V2 WHERE USER1_VF_S4556162.ATHLETE_V1.ATHLETEID = USER2_VF_S4556162.ATHLETE_V2.ATHLETEID AND USER1_VF_S4556162.ATHLETE_V1.ATHLETEID >= 305 AND USER1_VF_S4556162.ATHLETE_V1.ATHLETEID <= 310 AND USER2_VF_S4556162.ATHLETE_V2.ATHLETEID >= 305 AND USER2_VF_S4556162.ATHLETE_V2.ATHLETEID <= 310;		
SELECT FNAME, SNAME, BDATE FROM USER1_VF_S4556162.ATHLETE_V1, USER2_VF_S4556162.ATHLETE_V2 WHERE USER1_VF_S4556162.ATHLETE_V1.ATHLETEID = USER2_VF_S4556162.ATHLETE_V2.ATHLETEID AND USER1_VF_S4556162.ATHLETE_V1.ATHLETEID >= 305 AND USER1_VF_S4556162.ATHLETE_V1.ATHLETEID <= 310 AND USER2_VF_S4556162.ATHLETE_V2.ATHLETEID >= 305 AND USER2_VF_S4556162.ATHLETE_V2.ATHLETEID <= 310;		
FNAME	SNAME	BDATE
-----	-----	-----
"Jason"	"Kidd"	
"Chris"	"Paul"	
"Tayshaun"	"Prince"	
"Michael"	"Redd"	
"Dwyane"	"Wade"	
"Deron"	"Williams"	
6 rows selected		

Task 4

Semi Joins

STEP 1	SELECT DISTINCT ATHLETEID FROM USER1_VF_S4556162.ATHLETE_V1 A;		
	SELECT DISTINCT ATHLETEID FROM USER1_VF_S4556162.ATHLETE_V1 A;		
	First 10 Outputs		Final 10 Outputs
	ATHLETEID ----- 647 650 658 665 670 672 674 677 691 692	16474 16477 16479 16488 16512 16517 16521 16525 15895 15896 Only 5,000 rows currently supported in a script results 5,000 rows selected	
	NOTE: 5000 rows supported and displayed (in SQL Developer), actual output was 248584 rows (according to SQL Plus, specified in calculations below) 24584 rows selected.		
STEP 2	SELECT B.ATHLETEID, BDATE, CCODE, SPORTID FROM USER2_VF_S4556162.ATHLETE_V2 B WHERE B.ATHLETEID IN (SELECT DISTINCT A.ATHLETEID FROM USER1_VF_S4556162.ATHLETE_V1 A) AND B.CCODE = 'AUS';		
	SELECT B.ATHLETEID, BDATE, CCODE, SPORTID FROM USER2_VF_S4556162.ATHLETE_V2 B WHERE B.ATHLETEID IN (SELECT DISTINCT A.ATHLETEID FROM USER1_VF_S4556162.ATHLETE_V1 A) AND B.CCODE = 'AUS';		
	First 10 Outputs		
	ATHLETEID BDATE ----- 2188 2191 2193 2194 2195 2208 2214 2227 2228 2238	CCODE ----- AUS AUS AUS AUS AUS AUS AUS AUS AUS AUS	SPORTID ----- 11 11 11 11 11 11 11 11 11 12
	Final 10 Outputs		

	20324	AUS	38			
	20386	AUS	8			
	20390	AUS	8			
	22118	AUS	45			
	22139	AUS	45			
	22821	AUS	52			
	22858	AUS	45			
	23282	AUS	53			
	23283	AUS	53			
	23284	AUS	53			
	717 rows selected					
	Please Note: BDATE was visible for some of the other entries (besides the first/final 10 outputs) in which the field was specified. The below screenshots demonstrate some examples:					
	3 1982/10/13 0:00:00	AUS	33			
	11 1986/10/2 0:00:00	AUS	18			
	12	AUS	27			
	13 1971/1/25 0:00:00	AUS	11			
	14 1976/4/14 0:00:00	AUS	18			
	15 1964/2/24 0:00:00	AUS	11			
	These are consistent with the rest of the database's raw data provided.					
	NOTE: 717 rows displayed, actual output was 717 rows (according to both SQL Plus & SQL Developer, specified in calculations below)					
	717 rows selected.					
STEP 3	SELECT A.ATHLETEID, A.FNAME, A.SNAME, C.BDATE, C.CCODE, C.SPORTID FROM USER1_VF_S4556162.ATHLETE_V1 A, (SELECT B.ATHLETEID, B.BDATE, B.CCODE, B.SPORTID FROM USER2_VF_S4556162.ATHLETE_V2 B WHERE B.ATHLETEID IN (SELECT DISTINCT A.ATHLETEID FROM USER1_VF_S4556162.ATHLETE_V1 A) AND B.CCODE = 'AUS') C WHERE A.ATHLETEID = C.ATHLETEID;					
	SELECT A.ATHLETEID, A.FNAME, A.SNAME, C.BDATE, C.CCODE, C.SPORTID FROM USER1_VF_S4556162.ATHLETE_V1 A, (SELECT B.ATHLETEID, B.BDATE, B.CCODE, B.SPORTID FROM USER2_VF_S4556162.ATHLETE_V2 B WHERE B.ATHLETEID IN (SELECT DISTINCT A.ATHLETEID FROM USER1_VF_S4556162.ATHLETE_V1 A) AND B.CCODE = 'AUS') C WHERE A.ATHLETEID = C.ATHLETEID;					
	First 10 Outputs					
	ATHLETEID	FNAME	SNAME	BDATE	CCODE	SPORTID
	2188	"Ryan"	"Bayley"		AUS	11
	2191	"Brad"	"McGee"		AUS	11
	2193	"Graeme"	"Brown"		AUS	11
	2194	"Brett"	"Lancaster"		AUS	11
	2195	"Luke"	"Roberts"		AUS	11
	2208	"Sara"	"Carrigan"		AUS	11
2214	"Kate"	"Mactier"		AUS	11	
2227	"Shane"	"Kelly"		AUS	11	
2228	"Stuart"	"O'Grady"		AUS	11	
2238	"Mathew"	"Helm"		AUS	12	
	Final 10 Outputs					

11700 "Maureen"	"Caird"	AUS	38
11701 "Pamela"	"Kilborn"	AUS	38
22821 "Steven"	"Bradbury"	AUS	52
22858 "Zali"	"Steggall"	AUS	45
19389 "Nigel"	"Barker"	AUS	38
22118 "Dale"	"Begg-Smith"	AUS	45
22139 "Alisa"	"Campkin"	AUS	45
23282 "Andrew"	"Murtha"	AUS	53
23283 "Kieran"	"Hansen"	AUS	53
23284 "Richard"	"Nizielski"	AUS	53
1,434 rows selected			
Please Note: BDATE was visible for some of the other entries (besides the first/final 10 outputs) in which the field was specified. The below screenshots demonstrate some examples:			
3 "Ian"	"Thorpe"	1982/10/13 0:00:00	AUS 33
11 "Des"	"Abbott"	1986/10/2 0:00:00	AUS 18
12 "Michael"	"Aikman"		AUS 27
13 "Brett"	"Aitken"	1971/1/25 0:00:00	AUS 11
14 "Baeden"	"Choppy"	1976/4/14 0:00:00	AUS 18
15 "Michael"	"Grenda"	1964/2/24 0:00:00	AUS 11
These are consistent with the rest of the database's raw data provided.			
NOTE: 1434 rows displayed, actual output was 1434 rows (according to both SQL Plus & SQL Developer, specified in calculations below)			
1434 rows selected.			

Inner Joins

STEP 1	SELECT B.ATHLETEID, B.BDATE, B.CCODE, B.SPORTID FROM USER2_VF_S4556162.ATHLETE_V2 B WHERE B.CCODE = 'AUS';			
	SELECT B.ATHLETEID, B.BDATE, B.CCODE, B.SPORTID FROM USER2_VF_S4556162.ATHLETE_V2 B WHERE B.CCODE = 'AUS';			
	First 10 Outputs			
	ATHLETEID	BDATE	CCODE	SPORTID
	347		AUS	6
	348		AUS	6
	349		AUS	6
	350		AUS	6
	351		AUS	6
	352		AUS	6
	353		AUS	6
	354		AUS	6
	355		AUS	6
	356		AUS	6
	Final 10 Outputs			
	22821		AUS	52
	22858		AUS	45
	23282		AUS	53
	23283		AUS	53
	23284		AUS	53
	20212		AUS	33
	20324		AUS	38
	20386		AUS	8
	20390		AUS	8
	19389		AUS	38
717 rows selected				

STEP 2	NOTE: 717 rows displayed, actual output was 717 rows (according to both SQL Plus & SQL Developer, specified in calculations below) <div>717 rows selected.</div>				
	SELECT A.ATHLETEID, A.FNAME, A.SNAME, C.BDATE, C.CCODE, C.SPORTID FROM USER1_VF_S4556162.ATHLETE_V1 A, (SELECT B.ATHLETEID, B.BDATE, B.CCODE, B.SPORTID FROM USER2_VF_S4556162.ATHLETE_V2 B WHERE B.CCODE = 'AUS') C WHERE A.ATHLETEID = C.ATHLETEID;				
	SELECT A.ATHLETEID, A.FNAME, A.SNAME, C.BDATE, C.CCODE, C.SPORTID FROM USER1_VF_S4556162.ATHLETE_V1 A, (SELECT B.ATHLETEID, B.BDATE, B.CCODE, B.SPORTID FROM USER2_VF_S4556162.ATHLETE_V2 B WHERE B.CCODE = 'AUS') C WHERE A.ATHLETEID = C.ATHLETEID;				
	First 10 Outputs				
	ATHLETEID	FNAME	SNAME	BDATE	SPORTID
	2188	"Ryan"	"Bayley"		AUS 11
	2191	"Brad"	"McGee"		AUS 11
	2193	"Graeme"	"Brown"		AUS 11
	2194	"Brett"	"Lancaster"		AUS 11
	2195	"Luke"	"Roberts"		AUS 11
	2208	"Sara"	"Carriagan"		AUS 11
	2214	"Kate"	"Mactier"		AUS 11
	2227	"Shane"	"Kelly"		AUS 11
	2228	"Stuart"	"O'Grady"		AUS 11
	2238	"Mathew"	"Helm"		AUS 12
	Final 10 Outputs				
	11700	"Maureen"	"Caird"		AUS 38
	11701	"Pamela"	"Kilborn"		AUS 38
	22821	"Steven"	"Bradbury"		AUS 52
	22858	"Zali"	"Steggall"		AUS 45
	19389	"Nigel"	"Barker"		AUS 38
	22118	"Dale"	"Begg-Smith"		AUS 45
	22139	"Alisa"	"Campkin"		AUS 45
	23282	"Andrew"	"Murtha"		AUS 53
	23283	"Kieran"	"Hansen"		AUS 53
	23284	"Richard"	"Nizielski"		AUS 53
	1,434 rows selected				
	Please Note: BDATE was visible for some of the other entries (besides the first/final 10 outputs) in which the field was specified. The below screenshots demonstrate some examples:				
	3	"Ian"	"Thorpe"	1982/10/13 0:00:00	AUS 33
	11	"Des"	"Abbott"	1986/10/2 0:00:00	AUS 18
	12	"Michael"	"Aikman"		AUS 27
	13	"Brett"	"Aitken"	1971/1/25 0:00:00	AUS 11
	14	"Baeden"	"Choppy"	1976/4/14 0:00:00	AUS 18
	15	"Michael"	"Grenda"	1964/2/24 0:00:00	AUS 11
	These are consistent with the rest of the database's raw data provided.				
	NOTE: 1434 rows displayed, actual output was 1434 rows (according to both SQL Plus & SQL Developer, specified in calculations below) <div>1434 rows selected.</div>				

Global Join Query – Given in Practical 1 Task Sheet

select b.AthleteID, b.FName, b.SName, c.BDate, c.CCode, c.SportID from USER1_VF_S4556162.ATHLETE_V1 b, USER2_VF_S4556162.ATHLETE_V2 c where b.AthleteID= c.AthleteID and c.CCODE='AUS';				
select b.AthleteID, b.FName, b.SName, c.BDate, c.CCode, c.SportID from USER1_VF_S4556162.ATHLETE_V1 b, USER2_VF_S4556162.ATHLETE_V2 c where b.AthleteID= c.AthleteID and c.CCODE='AUS';				
First 10 Outputs				

ATHLETEID	FNAME	SNAME	BDATE	CCODE	SPORTID
2188	"Ryan"	"Bayley"		AUS	11
2191	"Brad"	"McGee"		AUS	11
2193	"Graeme"	"Brown"		AUS	11
2194	"Brett"	"Lancaster"		AUS	11
2195	"Luke"	"Roberts"		AUS	11
2208	"Sara"	"Carrigan"		AUS	11
2214	"Kate"	"Mactier"		AUS	11
2227	"Shane"	"Kelly"		AUS	11
2228	"Stuart"	"O'Grady"		AUS	11
2238	"Mathew"	"Helm"		AUS	12
2240	"Glen"	"Helm"		AUS	12
Final 10 Outputs					
11700	"Maureen"	"Caird"		AUS	38
11701	"Pamela"	"Kilborn"		AUS	38
22821	"Steven"	"Bradbury"		AUS	52
22858	"Zali"	"Steggall"		AUS	45
19389	"Nigel"	"Barker"		AUS	38
22118	"Dale"	"Begg-Smith"		AUS	45
22139	"Alisa"	"Camplin"		AUS	45
23282	"Andrew"	"Murtha"		AUS	53
23283	"Kieran"	"Hansen"		AUS	53
23284	"Richard"	"Nizielski"		AUS	53
1,434 rows selected					
NOTE: 1434 rows displayed, actual output was 1434 rows (according to both SQL Plus & SQL Developer, specified in calculations below)					
1434 rows selected.					

Calculations

Semi Joins

STEP 1	Time Elapsed: Elapsed: 00:00:22.69 Statistical Output: <pre> Statistics ----- 56 recursive calls 0 db block gets 287 consistent gets 0 physical reads 0 redo size 455587 bytes sent via SQL*Net to client 18570 bytes received via SQL*Net from client 1640 SQL*Net roundtrips to/from client 5 sorts (memory) 0 sorts (disk) 24584 rows processed </pre>
STEP 2	Time Elapsed: Elapsed: 00:00:00.86 Statistical Output:

	<pre> Statistics ----- 80 recursive calls 0 db block gets 314 consistent gets 0 physical reads 0 redo size 16949 bytes sent via SQL*Net to client 1069 bytes received via SQL*Net from client 49 SQL*Net roundtrips to/from client 13 sorts (memory) 0 sorts (disk) 717 rows processed </pre>
STEP 3	<p>Time Elapsed: Elapsed: 00:00:01.83</p> <p>Statistical Output:</p> <pre> Statistics ----- 3 recursive calls 0 db block gets 415 consistent gets 0 physical reads 0 redo size 59100 bytes sent via SQL*Net to client 1597 bytes received via SQL*Net from client 97 SQL*Net roundtrips to/from client 0 sorts (memory) 0 sorts (disk) 1434 rows processed </pre> <ul style="list-style-type: none"> NOTE: Key Factors are outlined in Red

Inner Joins

STEP 1	<p>Time Elapsed: Elapsed: 00:00:00.89</p> <p>Statistical Output:</p> <pre> Statistics ----- 1 recursive calls 0 db block gets 115 consistent gets 0 physical reads 0 redo size 16946 bytes sent via SQL*Net to client 1069 bytes received via SQL*Net from client 49 SQL*Net roundtrips to/from client 0 sorts (memory) 0 sorts (disk) 717 rows processed </pre>
STEP 2	<p>Time Elapsed: Elapsed: 00:00:01.76</p>

	<p>Statistical Output:</p> <pre> Statistics ----- 1 recursive calls 0 db block gets 409 consistent gets 0 physical reads 0 redo size 59100 bytes sent via SQL*Net to client 1597 bytes received via SQL*Net from client 97 SQL*Net roundtrips to/from client 0 sorts (memory) 0 sorts (disk) 1434 rows processed </pre>
	<ul style="list-style-type: none"> NOTE: Key Factors are outlined in Red

Global Join

<p>Time Elapsed:</p> <pre>Elapsed: 00:00:01.75</pre>	
	<p>Statistical Output:</p> <pre> Statistics ----- 1 recursive calls 0 db block gets 409 consistent gets 0 physical reads 0 redo size 59100 bytes sent via SQL*Net to client 1597 bytes received via SQL*Net from client 97 SQL*Net roundtrips to/from client 0 sorts (memory) 0 sorts (disk) 1434 rows processed </pre>
	<p>NOTE: Key Factors are outlined in Red</p>

In accordance with the task sheet, the method to calculate the transmission cost is as follows for both the semi-join and the inner-join plan. It is important to note that the most significant value for this calculation is the “bytes sent via SQL*NET to client” stated in the statistics. Hence, the calculations would be:

- Transmission Cost of the Semi Join Plan = 455 587 (step 1) + 16949 (step 2) = **472 536 bytes**
- Transmission Cost of the Inner Join Plan = **16 946 bytes** (step 1 only)

Therefore, the inner join plan is much better suited with regards to the global query, as the transmission cost is lower by 455 590 bytes in comparison to the semi join plan:

$$\begin{aligned} &\text{Transmission Cost of the Semi Join Plan} - \text{Transmission Cost of the Inner Join Plan} \\ &= 472\,536 - 16\,946 = \mathbf{455\,590\text{ bytes}} \end{aligned}$$

Similarly, the task sheet also states that the final step of the join execution plan, regardless of the method of join – either semi-join or inner-join in this case – needs to be equivalent. Hence, Step 3 of the semi-join plan, Step 2 of the inner-join plan and the global join should all match up – specifically in the statistics' categories outlined in red.

The following table compares and summaries the important statistics of the join execution plans for these three methods:

	Semi-Join Plan – Step 3	Inner-Join Plan – Step 2	Global Join Plan
Bytes sent via SQL*NET to client	59 100	59 100	59 100
Bytes received via SQL*NET to client	1597	1597	1597
SQL*NET roundtrips to/from client	97	97	97
Rows Processed	1434	1434	1434
Elapsed Time	00:00:01.83	00:00:01.76	00:00:01.75

It is evident that majority of the important statistics in the three join execution plans for the final step are the same and reflect the equivalence of the respective queries. The most discerning factor here would be that the elapsed time is faster for the inner join plan as compared with the semi join plan by about 7 milliseconds. This is primarily due to the use of the intermediate that the semi join plan has (Step 1, 2), while the inner join's Step 1 statistics are already very similar to semi join's Step 2 statistics.

Hence, it is evident that the inner join is the better suited plan, given the lower data transmission cost, the lower time elapsed and the use of less intermediates.

References

Introduction to PL/SQL Anonymous Block. (2020, April 11). Retrieved March 22, 2023, from <https://www.oracletutorial.com/plsql-tutorial/plsql-anonymous-block/>